

Floating Point

Adapted for CS367 @ GMU


Today: Floating Point

- Background: Fractional binary numbers
- IEEE floating point standard: Definition
- Example and properties
- Rounding, addition, multiplication
- Floating point in C
- Summary

Fractional binary numbers

- What is 1011.101_2 ?

Fractional Binary Numbers



■ Representation

- Bits to right of “binary point” represent fractional powers of 2
- Represents rational number:

$$\sum_{k=-j}^i b_k \times 2^k$$

Fractional Binary Numbers: Examples

Value	Representation
5 3/4	101.11 ₂
2 7/8	10.111 ₂
1 7/16	1.0111 ₂

Observations

- Divide by 2 by shifting right (unsigned)
- Multiply by 2 by shifting left
- Numbers of form 0.111111...₂ are just below 1.0
 - $1/2 + 1/4 + 1/8 + \dots + 1/2^i + \dots \rightarrow 1.0$
 - Use notation $1.0 - \varepsilon$

Representable Numbers

■ Limitation #1

- Can only exactly represent numbers of the form $x/2^k$
 - Other rational numbers have repeating bit representations
- Value Representation
 - $1/3$ $0.0101010101[01]..._2$
 - $1/5$ $0.001100110011[0011]..._2$
 - $1/10$ $0.0001100110011[0011]..._2$

■ Limitation #2

- Just one setting of binary point within the w bits
 - Limited range of numbers (very small values? very large?)

Today: Floating Point

- Background: Fractional binary numbers
- IEEE floating point standard: Definition
- Example and properties
- Rounding, addition, multiplication
- Floating point in C
- Summary

IEEE Floating Point

■ IEEE Standard 754

- Established in 1985 as uniform standard for floating point arithmetic
 - Before that, many idiosyncratic formats
- Supported by all major CPUs

■ Driven by numerical concerns

- Nice standards for rounding, overflow, underflow
- Hard to make fast in hardware
 - Numerical analysts predominated over hardware designers in defining standard

Floating Point Representation

■ Numerical Form:

$$(-1)^s M \cdot 2^E$$

- Sign bit **s** determines whether number is negative or positive
- Significand **M** normally a fractional value in range [1.0,2.0).
- Exponent **E** weights value by power of two

■ Encoding

- MSB **S** is sign bit **s**
- exp field encodes **E** (but is not equal to E)
- frac field encodes **M** (but is not equal to M)




Precision options


■ Single precision: 32 bits



■ Double precision: 64 bits



■ Extended precision: 80 bits (Intel only)



“Normalized” Values

$$v = (-1)^s M \cdot 2^E$$

- When: $\text{exp} \neq 000\ldots0$ and $\text{exp} \neq 111\ldots1$
- Exponent coded as a biased value: $E = \text{Exp} - \text{Bias}$
 - Exp: unsigned value of exp field
 - Bias = $2^{k-1} - 1$, where k is number of exponent bits
 - Single precision: 127 (Exp: 1...254, E: -126...127)
 - Double precision: 1023 (Exp: 1...2046, E: -1022...1023)
- Significand coded with implied leading 1: $M = 1.\text{xxx}\ldots\text{x}_2$
 - xxx...x: bits of frac field
 - Minimum when frac=000...0 ($M = 1.0$)
 - Maximum when frac=111...1 ($M = 2.0 - \epsilon$)
 - Get extra leading bit for “free”

Normalized Encoding Example

$$v = (-1)^s M \cdot 2^E$$

$$E = \text{Exp} - \text{Bias}$$

- Value: `float F = 15213.0;`
- $15213_{10} = 11101101101101_2$
 $= 1.1101101101101_2 \times 2^{13}$

- Significand
- $M = 1.\underline{1101101101101}_2$
- $\text{frac} = \underline{1101101101101}0000000000_2$

- Exponent
- $E = 13$
- $\text{Bias} = 127$
- $\text{Exp} = 140 = 10001100_2$

- Result:

0	10001100	110110110110100000000000
s	exp	frac

Practice Problem

The bias of the exponent is always $2^{(e-1)} - 1$, for e bits of exp.
What is the bias for each width of floating points?

- 8 bit fp (one of many "minifloat" formats), with 4 exp bits
- 16 bit fp ("binary16"), with 5 exp bits
- 32 bit fp ("single precision"), with 8 exp bits
- 64 bit fp ("double precision"), with 11 exp bits
- 128 bit fp ("quadruple precision"), with 15 exp bits
- Why do you think the exponent is usually a small portion of the total available bits?
 - What would happen if we had mostly exp and very little frac?

Practice Problem

Identify the values represented by these 8-bit floating point values. All are normalized. 1 sign bit, 4 exp bits, 3 frac bits.

$$\text{Bias} = 2^{(4-1)} - 1 = 8 - 1 = 7.$$

- 0 1001 010
- 1 0111 100
- 0 1110 000

Denormalized Values

$$v = (-1)^s M \cdot 2^E$$
$$E = 1 - \text{Bias}$$

- Condition: $\text{exp} = 000\dots0$
- Exponent value: $E = 1 - \text{Bias}$ (instead of $E = 0 - \text{Bias}$)
- Significand coded with implied leading 0: $M = 0.\text{xxx}\dots\text{x}_2$
 - $\text{xxx}\dots\text{x}$: bits of frac
- Cases
 - $\text{exp} = 000\dots0, \text{frac} = 000\dots0$
 - Represents zero value
 - Note distinct values: +0 and -0 (why?)
 - $\text{exp} = 000\dots0, \text{frac} \neq 000\dots0$
 - Numbers closest to 0.0
 - Equispaced

Practice Problem

Identify the values represented by these 8-bit floating point values. All are denormalized. 1 sign bit, 4 exp bits, 3 frac bits.
Bias = $2^{(4-1)-1} = 8-1=7$.

- 0 0000 010
- 1 0000 100
- 0 0000 000

Special Values


- Condition: $\text{exp} = 111\dots1$
- Case: $\text{exp} = 111\dots1, \text{frac} = 000\dots0$
 - Represents value ∞ (infinity)
 - Operation that overflows
 - Both positive and negative
 - E.g., $1.0/0.0 = -1.0/-0.0 = +\infty, 1.0/-0.0 = -\infty$
- Case: $\text{exp} = 111\dots1, \text{frac} \neq 000\dots0$
 - Not-a-Number (NaN)
 - Represents case when no numeric value can be determined
 - E.g., $\sqrt{-1}, \infty - \infty, \infty \times 0$

Practice Problem

■ Classify each of these 12-bit floating point numbers as normalized, denormalized, or special (and indicate what special value).

- 0 00101 110111
- 1 00000 111111
- 0 11111 000000
- 0 00000 100110
- 1 11111 111111
- 0 00000 000000
- 1 00001 101010


Visualization: Floating Point Encodings



Today: Floating Point

- Background: Fractional binary numbers
- IEEE floating point standard: Definition
- Example and properties
- Rounding, addition, multiplication
- Floating point in C
- Summary

Tiny Floating Point Example



■ 8-bit Floating Point Representation

- the sign bit is in the most significant bit
- the next four bits are the exponent, with a bias of 7
- the last three bits are the `frac`

■ Same general form as IEEE Format

- normalized, denormalized
- representation of 0, NaN, infinity

Dynamic Range (Positive Only)

	s	exp	frac	E	Value	
Denormalized numbers	0	0000	000	-6	0	
	0	0000	001	-6	$1/8*1/64 = 1/512$	
	0	0000	010	-6	$2/8*1/64 = 2/512$	
	...					
	0	0000	110	-6	$6/8*1/64 = 6/512$	
	0	0000	111	-6	$7/8*1/64 = 7/512$	
	0	0001	000	-6	$8/8*1/64 = 8/512$	largest denorm
Normalized numbers	0	0001	001	-6	$9/8*1/64 = 9/512$	smallest norm
	...					
	0	0110	110	-1	$14/8*1/2 = 14/16$	
	0	0110	111	-1	$15/8*1/2 = 15/16$	closest to 1 below
	0	0111	000	0	$8/8*1 = 1$	
	0	0111	001	0	$9/8*1 = 9/8$	closest to 1 above
	0	0111	010	0	$10/8*1 = 10/8$	
	...					
	0	1110	110	7	$14/8*128 = 224$	
	0	1110	111	7	$15/8*128 = 240$	largest norm
	0	1111	000	n/a	inf	

$$v = (-1)^s M \cdot 2^E$$

n: E = Exp – Bias
d: E = 1 – Bias


closest to zero

largest denorm
smallest norm

closest to 1 below
closest to 1 above

largest norm

Examples: binary to float



1. 0 0101 110

$$\text{exp} = 5, \text{ so } E = 5 - 7 = -2$$

$$\text{frac} = 6/8, \text{ so } M = 14/8 = 14/8 * 2^{-2} = 7/16$$

2. 1 1010 011

$$\text{exp} = 10, \text{ so } E = 10 - 7 = 3$$

$$\text{frac} = 3/8, \text{ so } M = 11/8 = -11/8 * 2^3 = -11.0$$


3. 0 0000 010

$\text{exp} = 0$, so denormalized.

$$E = -6, (-\text{bias} + 1)$$


$$\text{frac} = 2/8, \text{ so } M = 2/8 = 2/8 * 1/64 = 1/256$$

Examples: float to binary



1. $52.0 = 1.625 * 2^5$
 - E = 5 so exp must be 12
 - $1.625 = 13/8$ so $\text{frac} = 5/8 = 0\ 1100\ 101$
 - Could also do this way: 52 is 110100 in 6-bits.
 - Dividing by 2 5 times gives us 1.10100

2. $7/64 = 14/8 * 1/2^4$
 - E = -4 so exp must be 3
 - $\text{frac} = 6/8 = 0\ 0011\ 110$



Practice Problem

- What is the largest denormalized value for 8-bit numbers? (bit pattern, value)
- What is the smallest normalized value for 8-bit numbers? (bit pattern, value)

Practice Problem


Write each of these numbers in base two scientific notation. Then, represent them as 8-bit floating point values.

- a. 3
- b. 32
- c. - 40
- d. 100
- e. $3/512$
- f. $-5/64$
- g. 1


Distribution of Values

■ 6-bit IEEE-like format

- $e = 3$ exponent bits
- $f = 2$ fraction bits
- Bias is $2^{3-1}-1 = 3$





■ Notice how the distribution gets denser toward zero.



Distribution of Values (close-up view)

■ 6-bit IEEE-like format

- $e = 3$ exponent bits
- $f = 2$ fraction bits
- Bias is 3



Special Properties of the IEEE Encoding

- FP Zero Same as Integer Zero
 - All bits = 0
- Can (Almost) Use Unsigned Integer Comparison
 - Must first compare sign bits
 - Must consider $-0 = 0$
 - NaNs problematic
 - Will be greater than any other values
 - What should comparison yield?
 - Otherwise OK
 - Denorm vs. normalized
 - Normalized vs. infinity

Today: Floating Point

- Background: Fractional binary numbers
- IEEE floating point standard: Definition
- Example and properties
- Rounding, addition, multiplication
- Floating point in C
- Summary

Floating Point Operations: Basic Idea

- $x +_f y = \text{Round}(x + y)$
- $x \times_f y = \text{Round}(x \times y)$
- Basic idea
 - First **compute exact result**
 - Make it fit into desired precision
 - Possibly overflow if exponent too large
 - Possibly **round to fit into `frac`**

Rounding

■ Rounding Modes (illustrate with \$ rounding)

	\$1.40	\$1.60	\$1.50	\$2.50	-\$1.50
■ Towards zero	\$1	\$1	\$1	\$2	-\$1
■ Round down ($-\infty$)	\$1	\$1	\$1	\$2	-\$2
■ Round up ($+\infty$)	\$2	\$2	\$2	\$3	-\$1
■ Nearest Even (default)	\$1	\$2	\$2	\$2	-\$2

Closer Look at Round-To-Even

■ Default Rounding Mode

- Hard to get any other kind without dropping into assembly
- All others are statistically biased
 - Sum of set of positive numbers will consistently be over- or under-estimated

■ Applying to Other Decimal Places / Bit Positions

- When exactly halfway between two possible values
 - Round so that least significant digit is even
- E.g., round to nearest hundredth

7.8949999	7.89	(Less than half way)
-----------	------	----------------------

7.8950001	7.90	(Greater than half way)
-----------	------	-------------------------

7.8950000	7.90	(Half way—round up)
-----------	------	---------------------

7.8850000	7.88	(Half way—round down)
-----------	------	-----------------------

Rounding Binary Numbers

■ Binary Fractional Numbers

- “Even” when least significant bit is 0
- “Half way” when bits to right of rounding position = $100\dots_2$

■ Examples

- Round to nearest 1/4 (2 bits right of binary point)

Value	Binary	Rounded	Action	Rounded Value
2 3/32	10.00 011 $_2$	10.00 $_2$	(<1/2—down)	2
2 3/16	10.00 110 $_2$	10.01 $_2$	(>1/2—up)	2 1/4
2 7/8	10.11 100 $_2$	11.00 $_2$	(1/2—up)	3
2 5/8	10.10 100 $_2$	10.10 $_2$	(1/2—down)	2 1/2

FP Multiplication

- $(-1)^{s_1} M_1 2^{E_1} \times (-1)^{s_2} M_2 2^{E_2}$
- Exact Result: $(-1)^s M 2^E$
 - Sign s: $s_1 \wedge s_2$
 - Significand M: $M_1 \times M_2$
 - Exponent E: $E_1 + E_2$
- Fixing
 - If $M \geq 2$, shift M right, increment E
 - If E out of range, overflow
 - Round M to fit `frac` precision
- Implementation
 - Biggest chore is multiplying significands

Floating Point Addition

- $(-1)^{s_1} M_1 2^{E_1} + (-1)^{s_2} M_2 2^{E_2}$


- Assume $E_1 > E_2$

- Exact Result: $(-1)^s M 2^E$

- Sign s, significand M:
 - Result of signed align & add

- Exponent E: E_1

Get binary points lined up



Fixing

- If $M \geq 2$, shift M right, increment E
- If $M < 1$, shift M left k positions, decrement E by k
- Overflow if E out of range
- Round M to fit $\frac{1}{2}$ precision

Mathematical Properties of FP Add

■ Compare to those of Abelian Group

- Closed under addition?
 - But may generate infinity or NaN
- Commutative?
Yes
- Associative?
 - Overflow and inexactness of rounding
 - $(3.14 + 1e10) - 1e10 = 0, 3.14 + (1e10 - 1e10) = 3.14$
- 0 is additive identity?
Yes
- Every element has additive inverse?
 - Yes, except for infinities & NaNs
Almost

■ Monotonicity

- $a \geq b \Rightarrow a+c \geq b+c?$
 - Except for infinities & NaNs
Almost

Mathematical Properties of FP Mult

■ Compare to Commutative Ring

- Closed under multiplication? Yes
 - But may generate infinity or NaN
- Multiplication Commutative? Yes
- Multiplication is Associative? No
 - Possibility of overflow, inexactness of rounding
 - Ex: $(1e20 * 1e20) * 1e-20 = \inf$, $1e20 * (1e20 * 1e-20) = 1e20$
- 1 is multiplicative identity? Yes
- Multiplication distributes over addition? No
 - Possibility of overflow, inexactness of rounding
 - $1e20 * (1e20 - 1e20) = 0.0$, $1e20 * 1e20 - 1e20 * 1e20 = \text{NaN}$

■ Monotonicity

- $a \geq b \ \& \ c \geq 0 \Rightarrow a * c \geq b * c?$ Almost
 - Except for infinities & NaNs

Today: Floating Point

- Background: Fractional binary numbers
- IEEE floating point standard: Definition
- Example and properties
- Rounding, addition, multiplication
- Floating point in C
- Summary

Floating Point in C

■ C Guarantees Two Levels

- float single precision
- double double precision

■ Conversions/Casting

- Casting between int, float, and double changes bit representation
- double/float → int
 - Truncates fractional part
 - Like rounding toward zero
 - Not defined when out of range or NaN: Generally sets to TMin
- int → double
 - Exact conversion, as long as int has \leq 53 bit word size
- int → float
 - Will round according to rounding mode

Floating Point Puzzles

- For each of the following C expressions, either:

- Argue that it is true for all argument values
- Explain why not true

```
int x = ...;
float f = ...;
double d = ...;
```

Assume neither
d nor f is NaN

- a) $x == (\text{int})(\text{float}) \ x$
- b) $x == (\text{int})(\text{double}) \ x$
- c) $f == (\text{float})(\text{double}) \ f$
- d) $d == (\text{double})(\text{float}) \ d$
- e) $f == -(-f);$
- f) $2/3 == 2/3.0$
- g) $d < 0.0 \Rightarrow ((d*2) < 0.0)$
- h) $d > f \Rightarrow -f > -d$
- i) $d * d \geq 0.0$
- j) $(d+f)-d == f$


Summary

- IEEE Floating Point has clear mathematical properties
- Represents numbers of form $M \times 2^E$
- One can reason about operations independent of implementation
 - As if computed with perfect precision and then rounded
- Not the same as real arithmetic
 - Violates associativity/distributivity
 - Makes life difficult for compilers & serious numerical applications programmers

Creating Floating Point Number

■ Steps

- Normalize to have leading 1
- Round to fit within fraction
- Postnormalize to deal with effects of rounding




■ Case Study

- Convert 8-bit unsigned numbers to tiny floating point format

Example Numbers

128	10000000
13	00001101
33	00010001
35	00010011
138	10001010
63	00111111

Normalize




Requirement

- Set binary point so that numbers of form 1.xxxxxx
- Adjust all to have leading one
 - Decrement exponent as shift left

Value	Binary	Fraction	Exponent
128	10000000	1.0000000	7
13	00001101	1.1010000	3
17	00010001	1.0001000	4
19	00010011	1.0011000	4
138	10001010	1.0001010	7
63	00111111	1.1111100	5

Rounding



■ Round up conditions

- Round = 1, Sticky = 1 $\rightarrow > 0.5$
- Guard = 1, Round = 1, Sticky = 0 \rightarrow Round to even

Value	Fraction	GRS	Incr?	Rounded
128	1.000 0000	000	N	1.000
13	1.101 0000	100	N	1.101
17	1.000 1000	010	N	1.000
19	1.001 1000	110	Y	1.010
138	1.000 1010	011	Y	1.001
63	1.111 1100	111	Y	10.000

Postnormalize

■ Issue

- Rounding may have caused overflow
- Handle by shifting right once & incrementing exponent

Value	Rounded	Exp	Adjusted	Result
128	1.000	7		128
13	1.101	3		13
17	1.000	4		16
19	1.010	4		20
138	1.001	7		144
63	10.000	5	1.000/6	64

Interesting Numbers

{single, double}

<i>Description</i>	<i>exp</i>	<i>frac</i>	<i>Numeric Value</i>
■ Zero	00...00	00...00	0.0
■ Smallest Pos. Denorm.	00...00	00...01	$2^{-\{23,52\}} \times 2^{-\{126,1022\}}$
■ Single $\approx 1.4 \times 10^{-45}$			
■ Double $\approx 4.9 \times 10^{-324}$			
■ Largest Denormalized	00...00	11...11	$(1.0 - \varepsilon) \times 2^{-\{126,1022\}}$
■ Single $\approx 1.18 \times 10^{-38}$			
■ Double $\approx 2.2 \times 10^{-308}$			
■ Smallest Pos. Normalized	00...01	00...00	$1.0 \times 2^{-\{126,1022\}}$
■ Just larger than largest denormalized			
■ One	01...11	00...00	1.0
■ Largest Normalized	11...10	11...11	$(2.0 - \varepsilon) \times 2^{\{127,1023\}}$
■ Single $\approx 3.4 \times 10^{38}$			
■ Double $\approx 1.8 \times 10^{308}$			

Ariane 5

- Exploded 37 seconds after liftoff
- Cargo worth \$500 million

Why did this happen?

- Computed horizontal velocity as floating point number
- Converted to 16-bit integer
- Worked OK for Ariane 4
- Overflowed for the more powerful Ariane 5

Used same software without re-checking assumptions

