



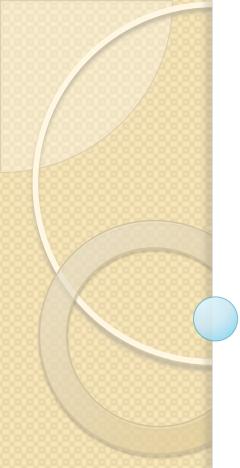
Final QuickFire Review

Fall 2012

CS 465

Huzefa Rangwala, PhD

Chapters 1 - 6



Disclaimer:

These slides are not a substitute for reading the book.

To do well: Go over the detailed readings, solve exercises from HW, Quizzes, Classroom Exercises and back of the book*.



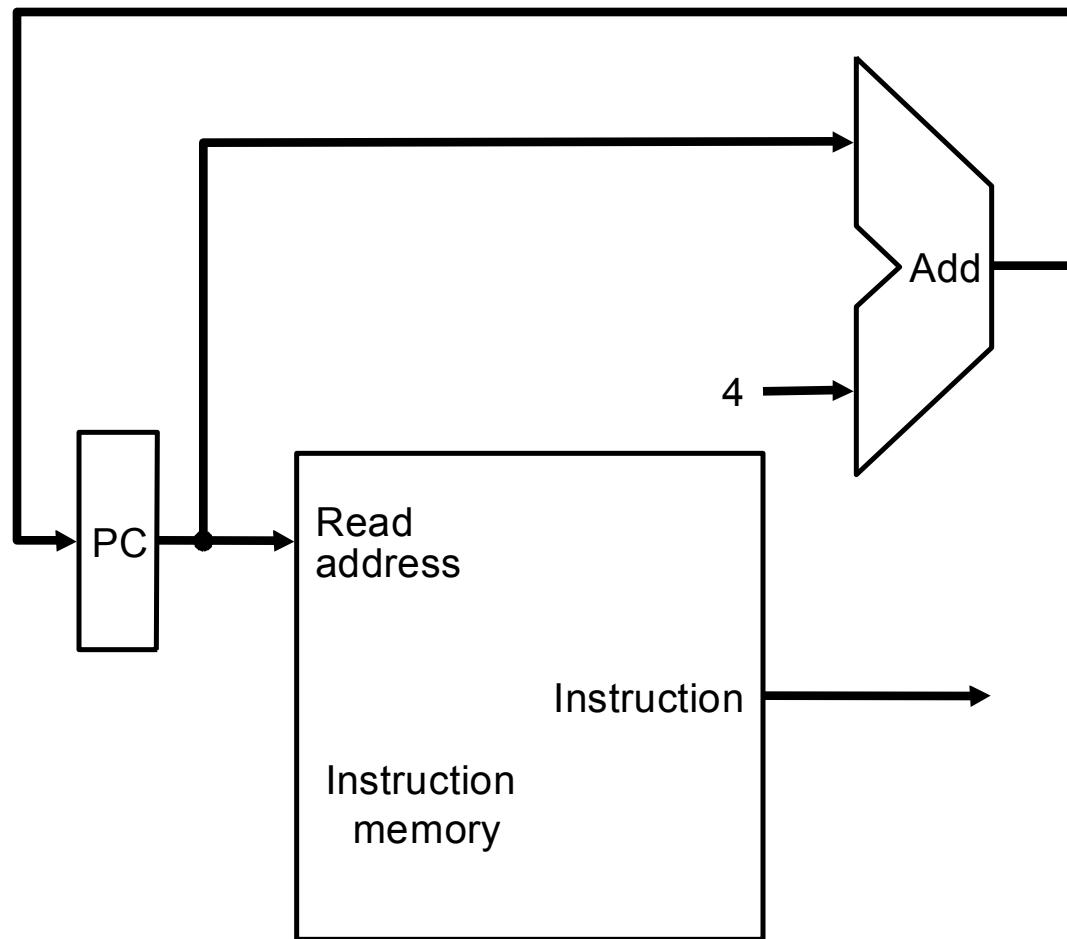
Detailed Readings

- Look on piazza or class web-site

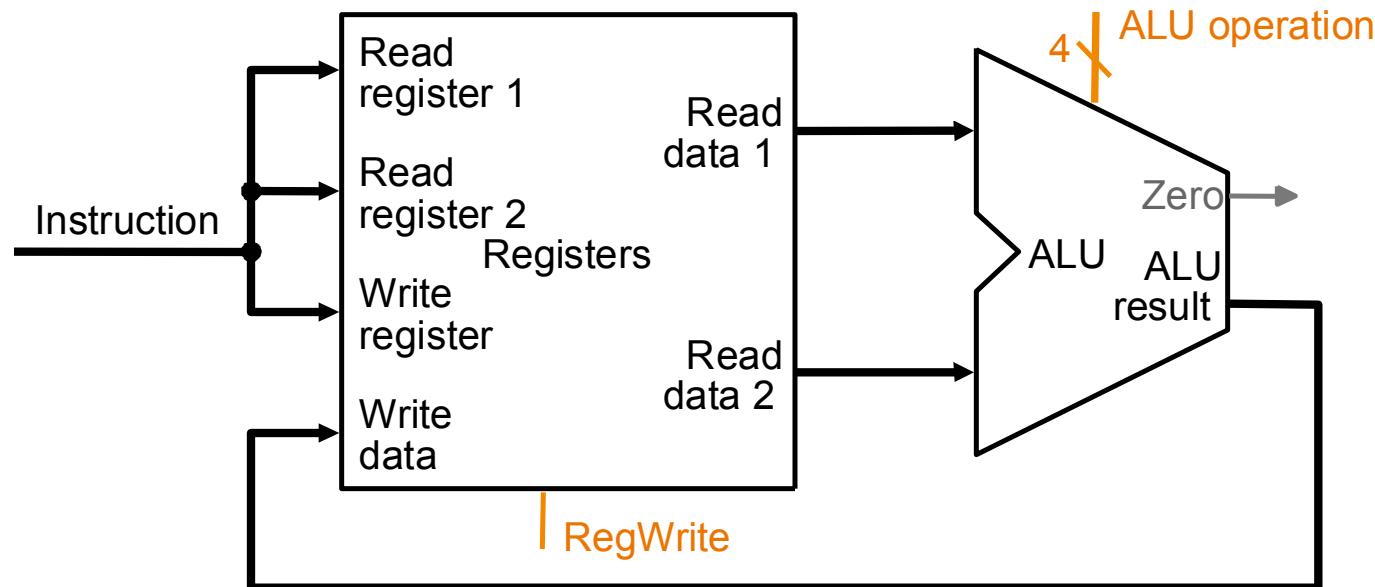
How to Design a Processor

- Analyze instruction set \Rightarrow datapath requirements
 - Meaning of each instruction given by register transfers
 - Datapath must include storage element for ISA registers (possibly more)
 - Datapath must support each register transfer
- Select set of datapath components and establish clocking methodology
- Assemble datapath meeting the requirements
- Analyze implementation of each instruction to determine setting of control points that effects the register transfer
- Assemble the control logic

Datapath for Instruction Fetch

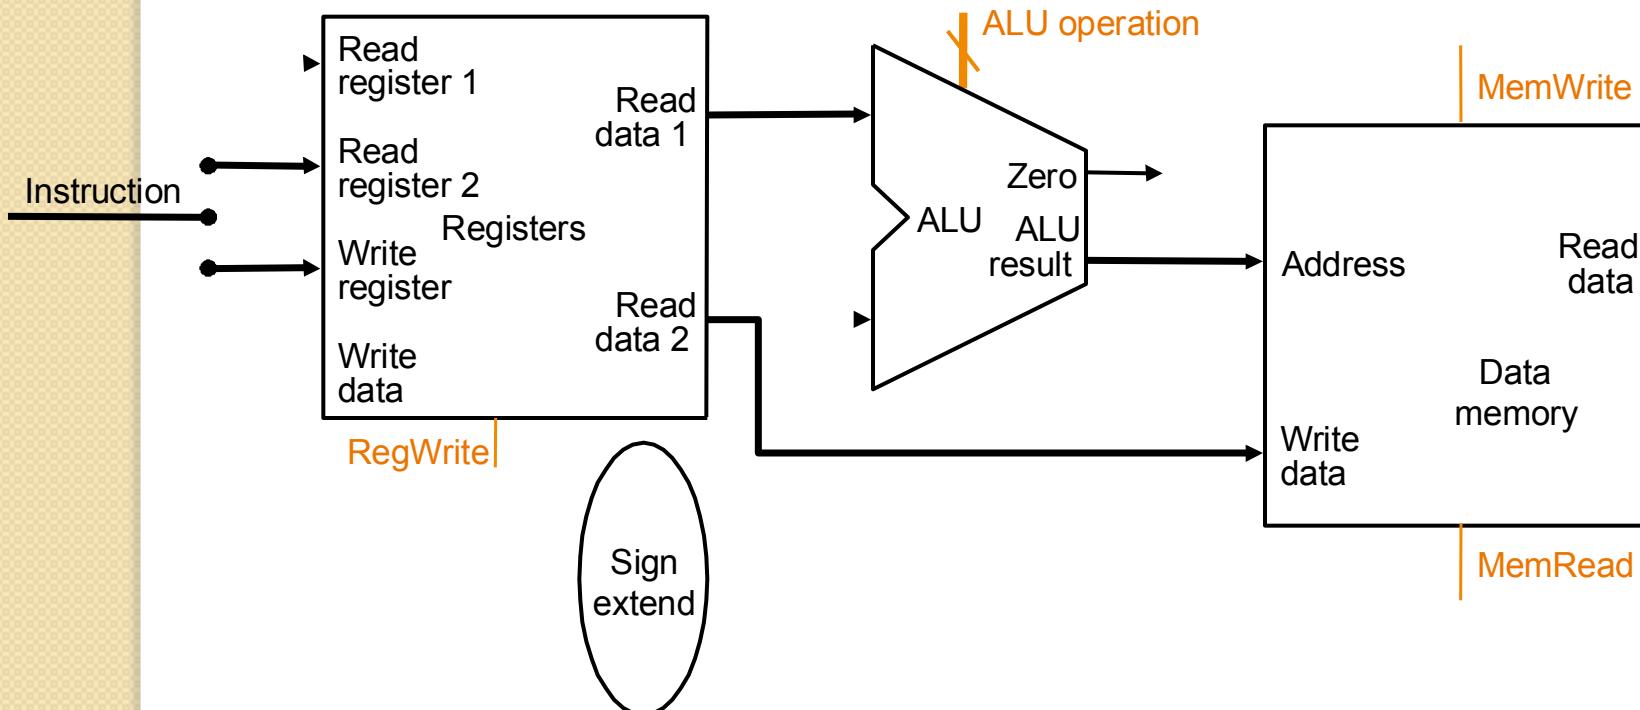


Datapath for R-format Instructions

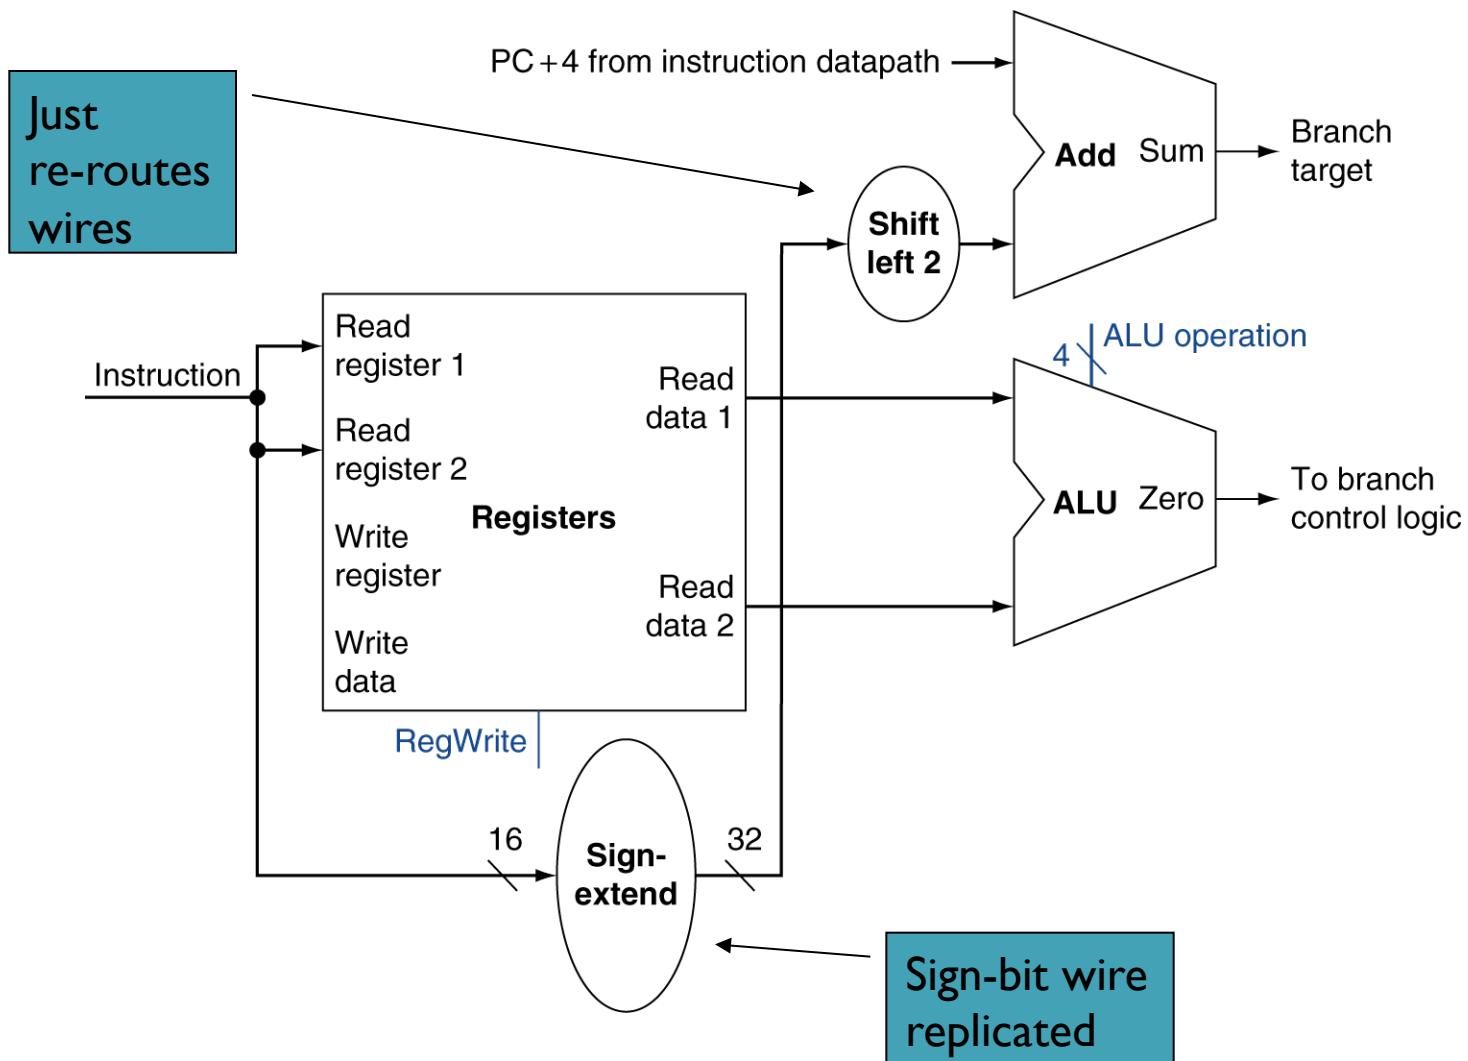


Datapath for LW & SW (Fix this)

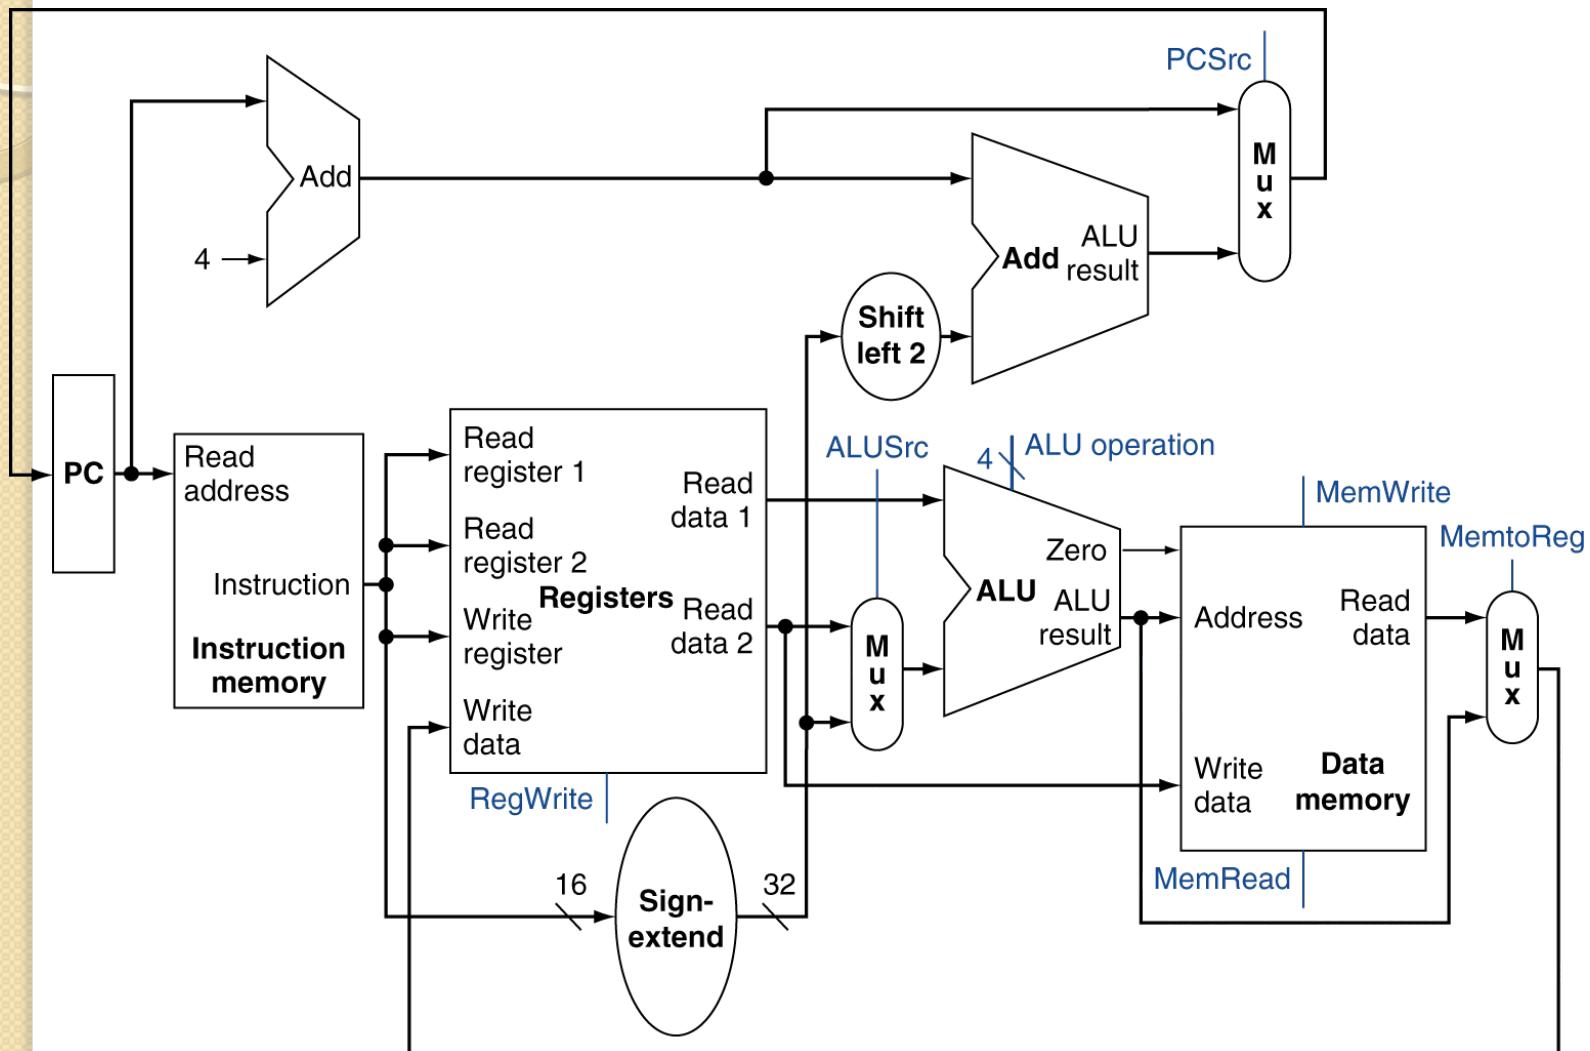
- ❑ $R[rt] \leftarrow \text{Mem}[R[rs] + \text{SignExt}[\text{imm16}]]$ Example: lw rt, rs, imm16
- ❑ $\text{Mem}[R[rs] + \text{SignExt}[\text{imm16}]] \leftarrow R[rt]$ Example: sw rt, rs, imm16



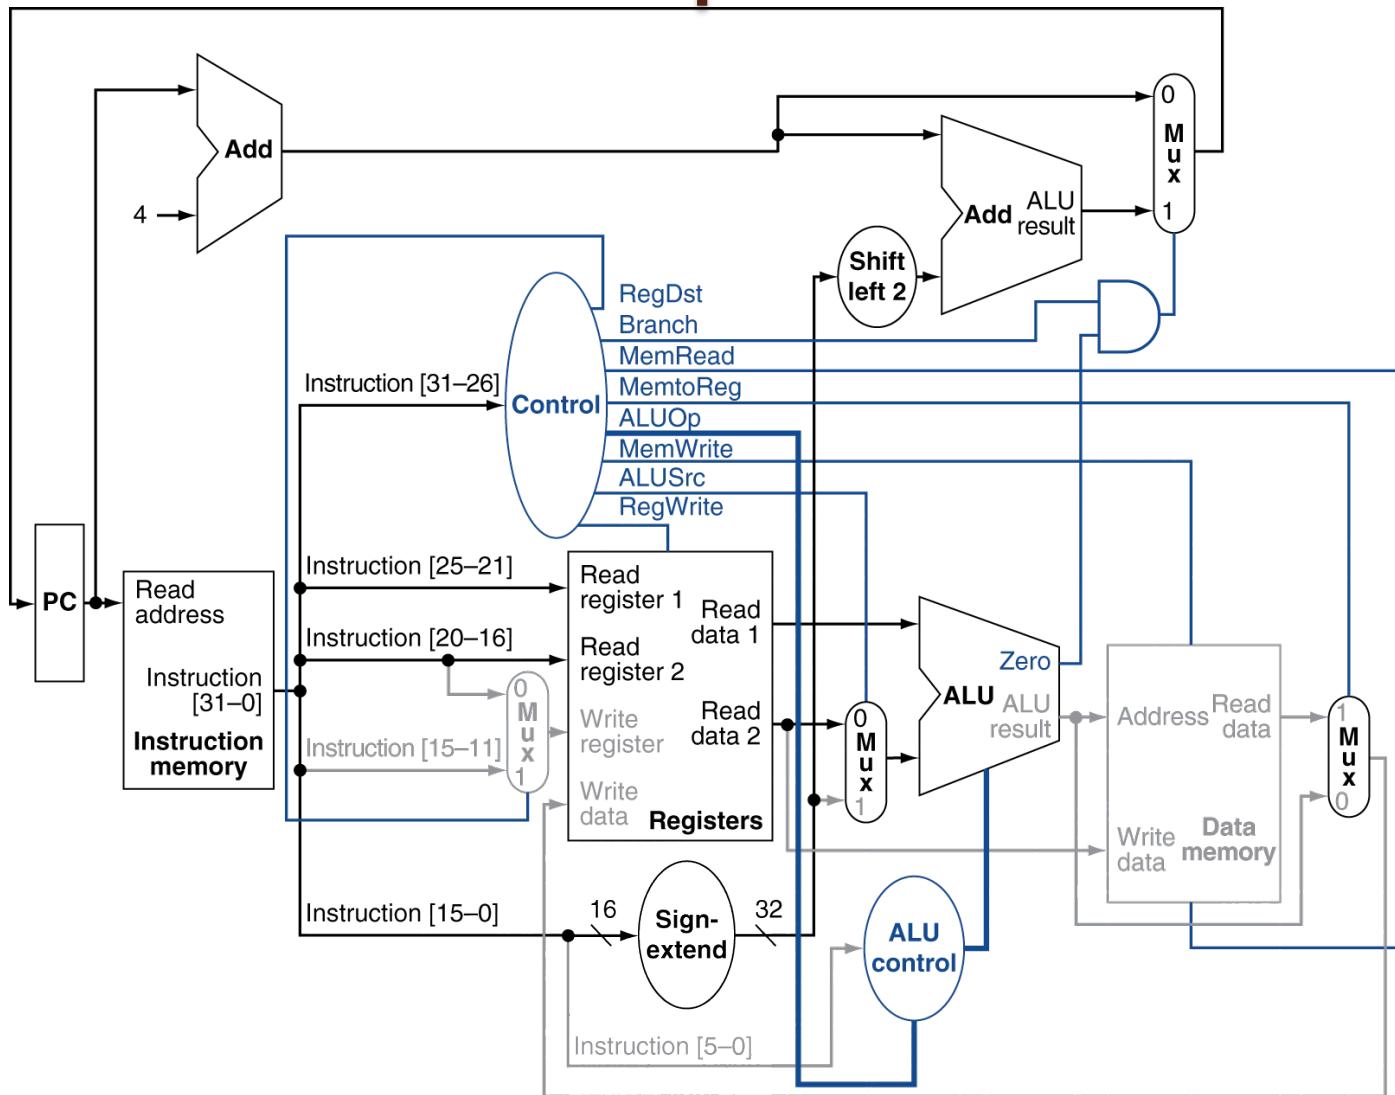
Branch Instructions



Full Datapath

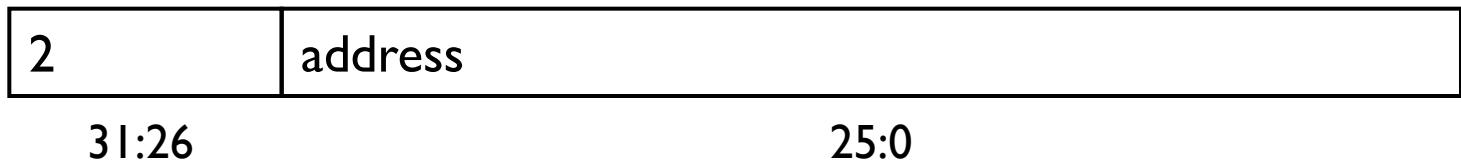


Branch-on-Equal Instruction



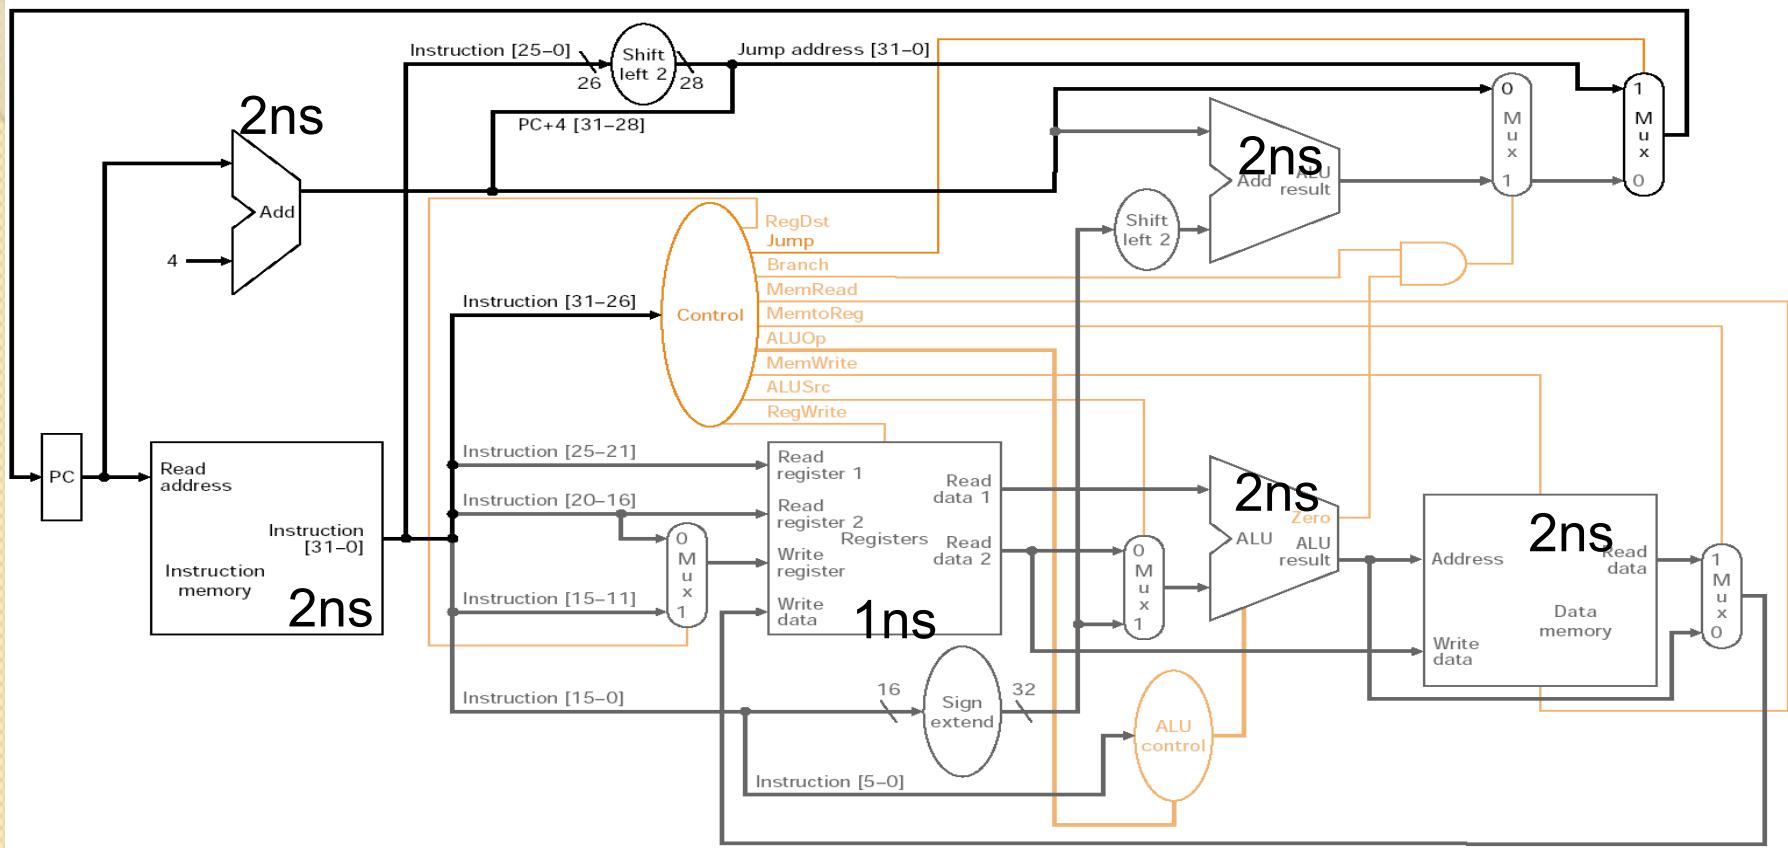
Implementing Jumps

Jump



- Jump uses word address
- Update PC with concatenation of
 - Top 4 bits of old PC
 - 26-bit jump address
 - 00
- Need an extra control signal decoded from opcode

Delays in Single Cycle Datapath



What are the delays for **lw, sw, R-Type, beq, j** instructions?

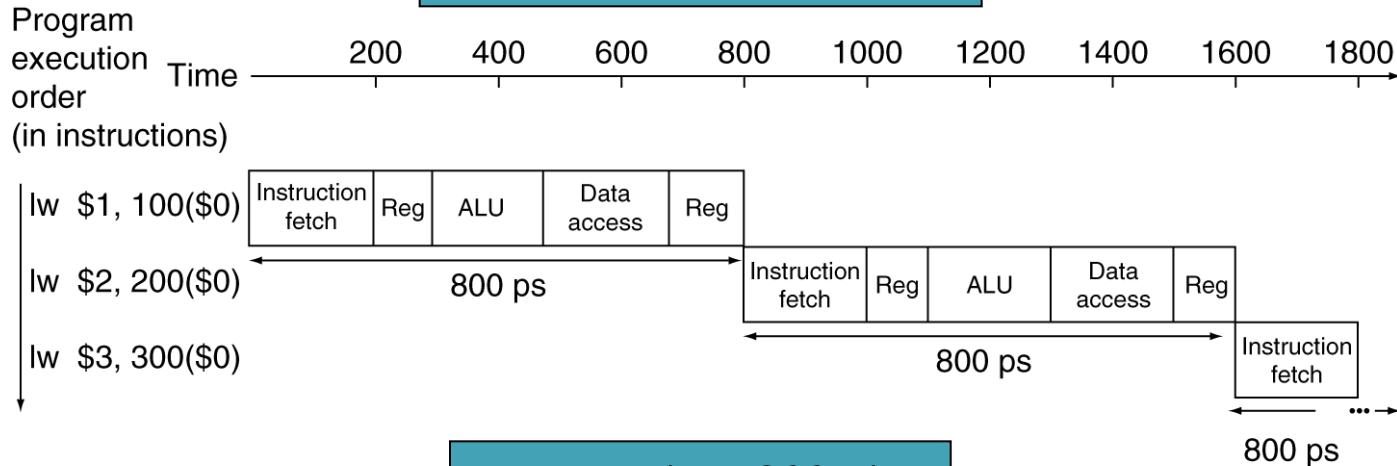
Single Cycle Implementation

- Calculate cycle time assuming negligible delays except:
 - memory (2ns), ALU and adders (2ns), register file access (1ns)

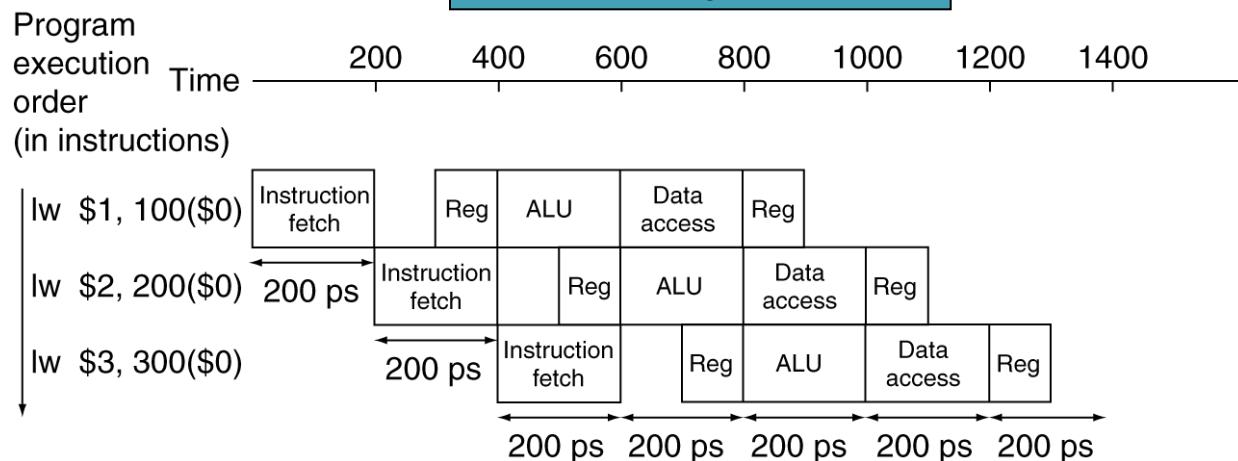
Instruction class	Instruction Fetch	Register Access	ALU	Register/ Memory Access	Register Access	
R-Type	X	X	X	R		6
Load	X	X	X	M	X	8
Store	X	X	X	M		7
Branch	X	X	X			5
Jump	X					2

Pipeline Performance

Single-cycle ($T_c = 800\text{ps}$)



Pipelined ($T_c = 200\text{ps}$)

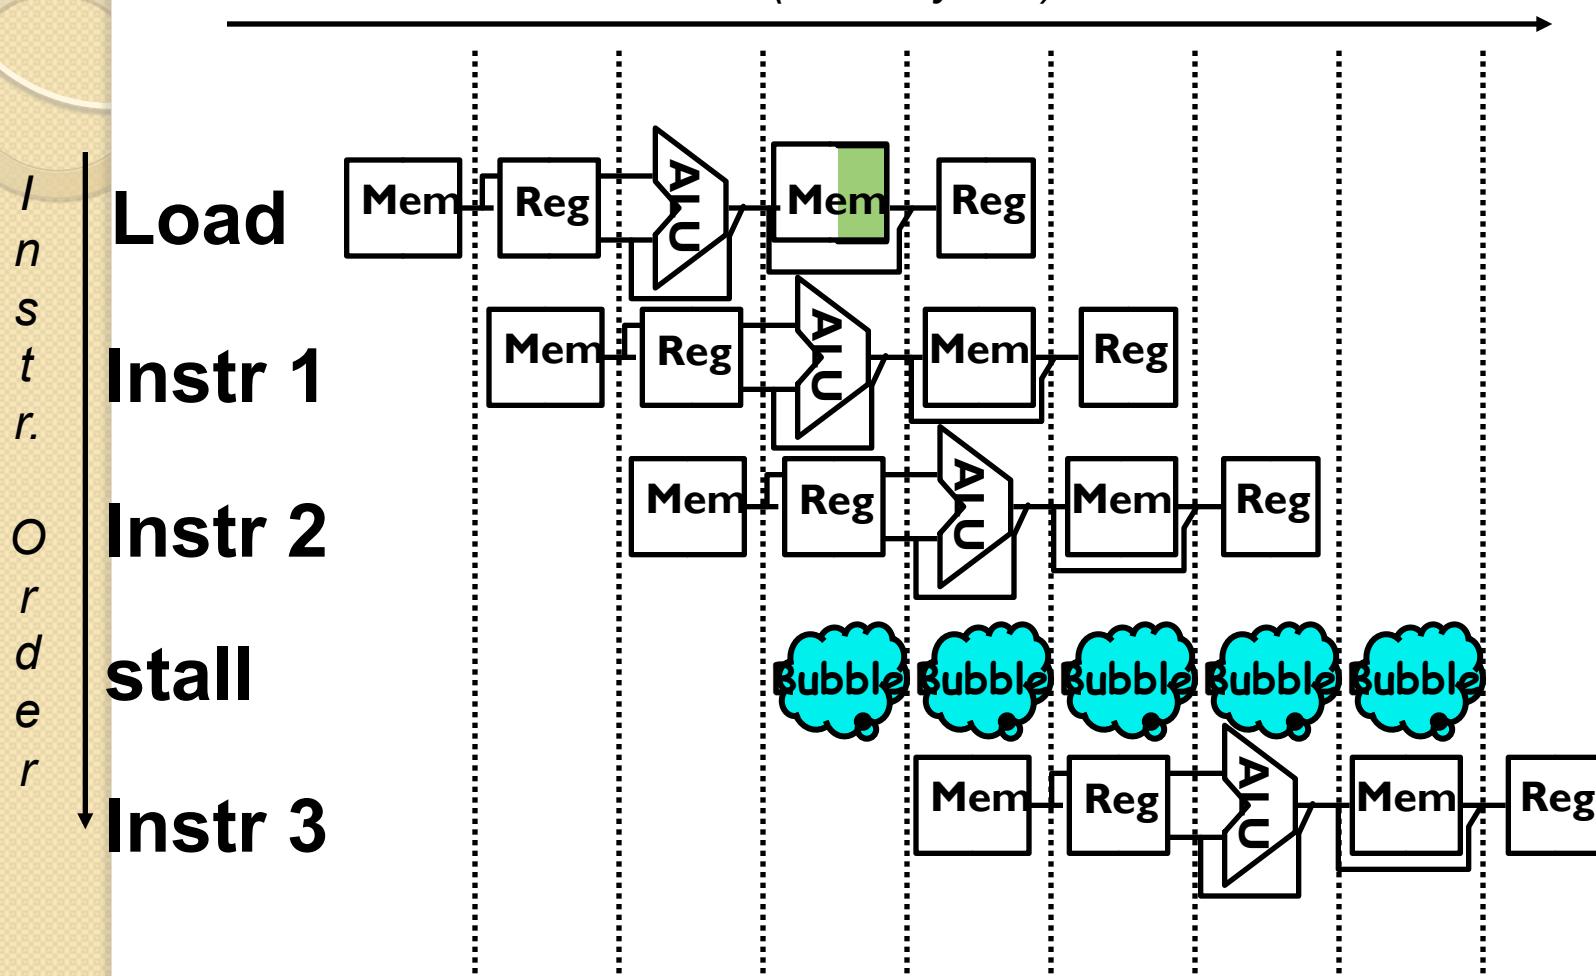


Hazards

- Situations that prevent starting the next instruction in the next cycle
- _____ hazards
 - A required resource is busy
- _____ hazard
 - Need to wait for previous instruction to complete its data read/write
- _____ hazard
 - Deciding on control action depends on previous instruction

Structural Hazard: One Memory

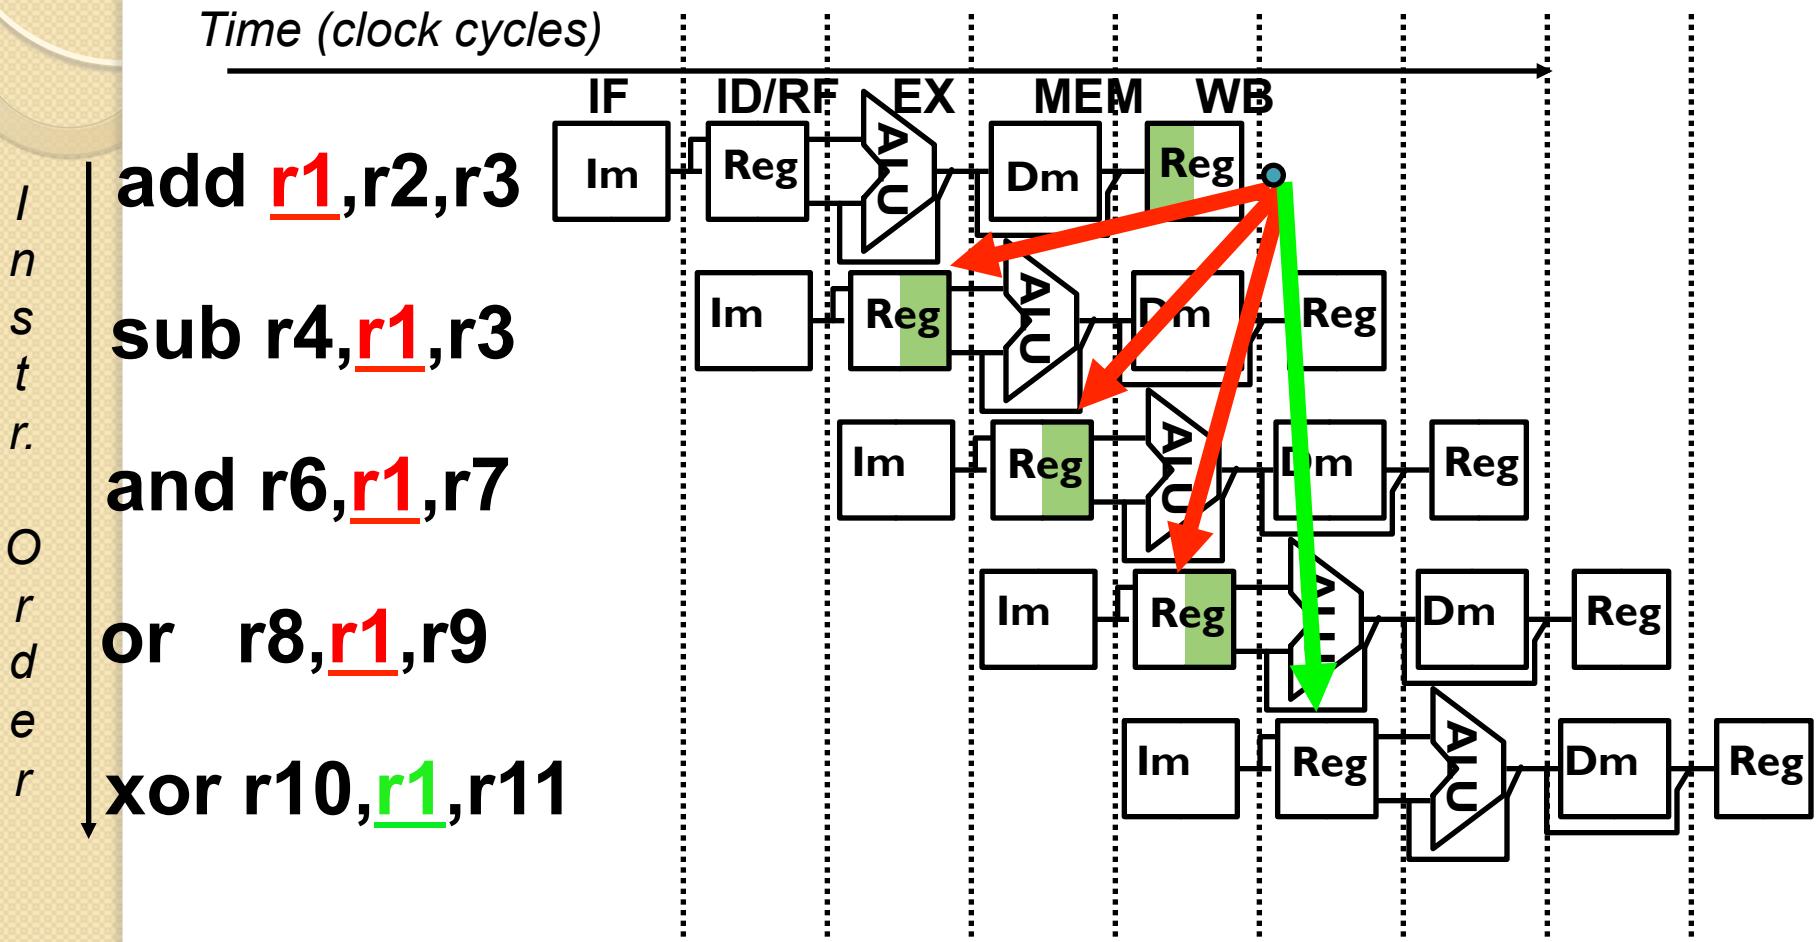
Time (clock cycles)



- Hazards can always be resolved by waiting

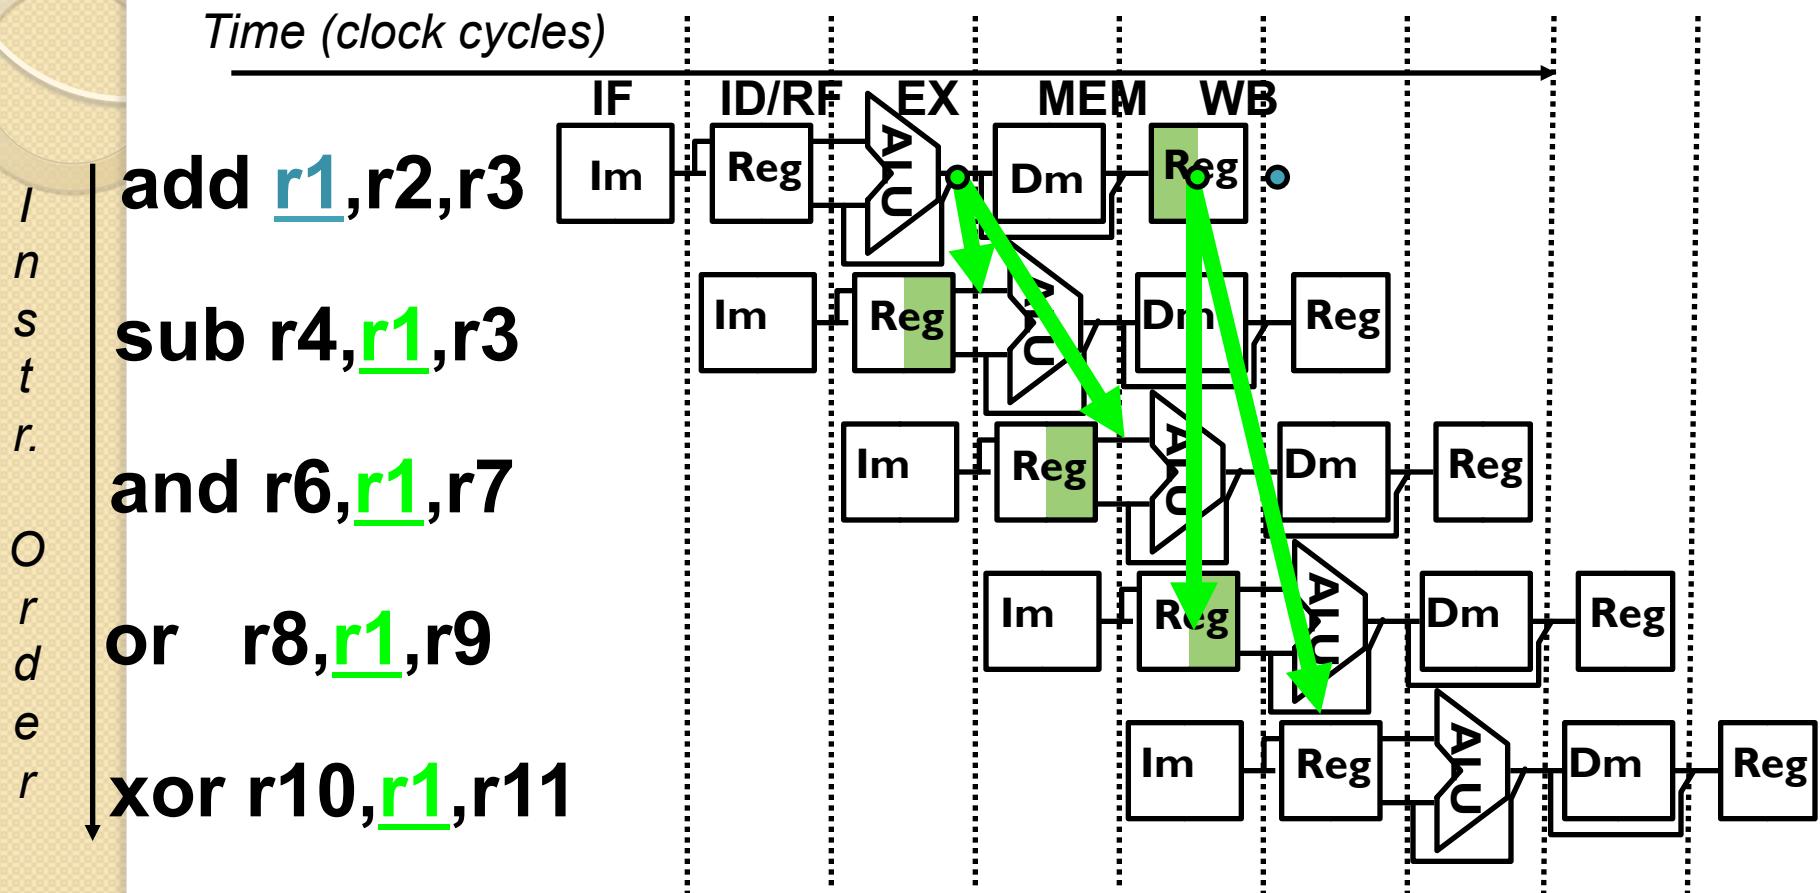
Data Hazard Example

- Dependences backward in time are hazards



- Compilers can help, but it gets messy and difficult

Data Hazard Solution



- Solution : “forward” result from one stage to another

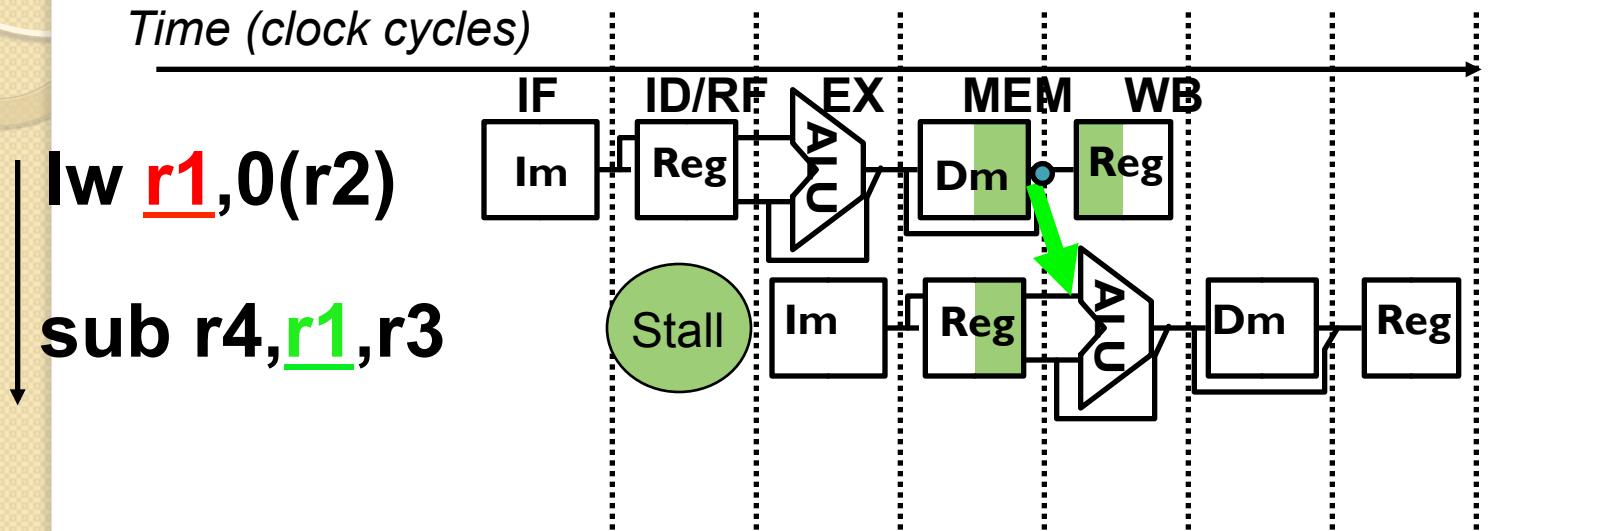
Data Hazard Even with Forwarding

- Can you resolve this ?

lw r1,0(r2)

sub r4,r1,r3

Data Hazard Even with Forwarding



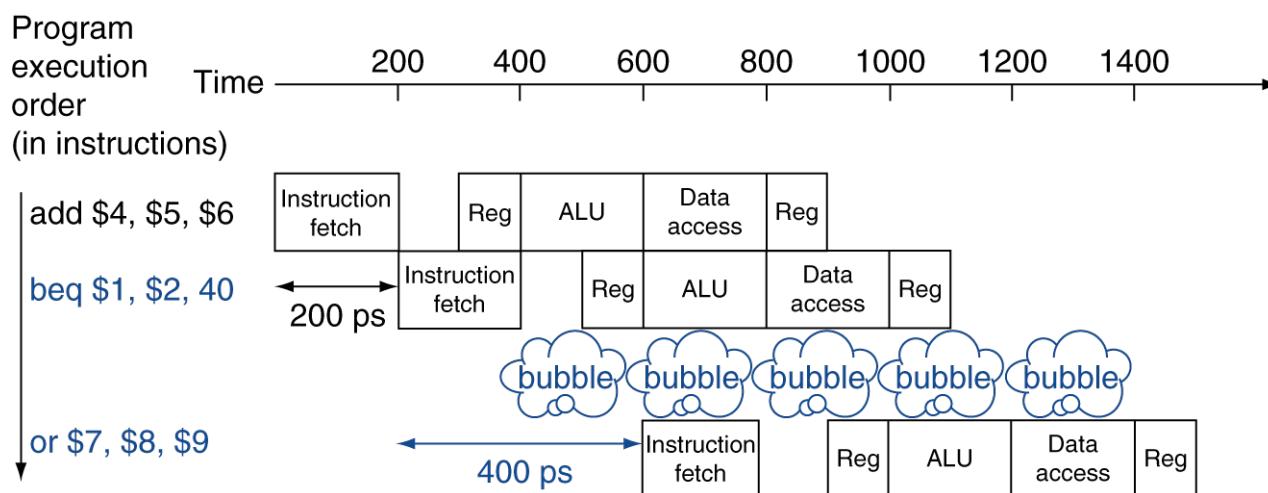
- Must delay/stall instruction dependent on loads
- Sometimes the instruction sequence can be reordered to avoid pipeline stalls

Control Hazards

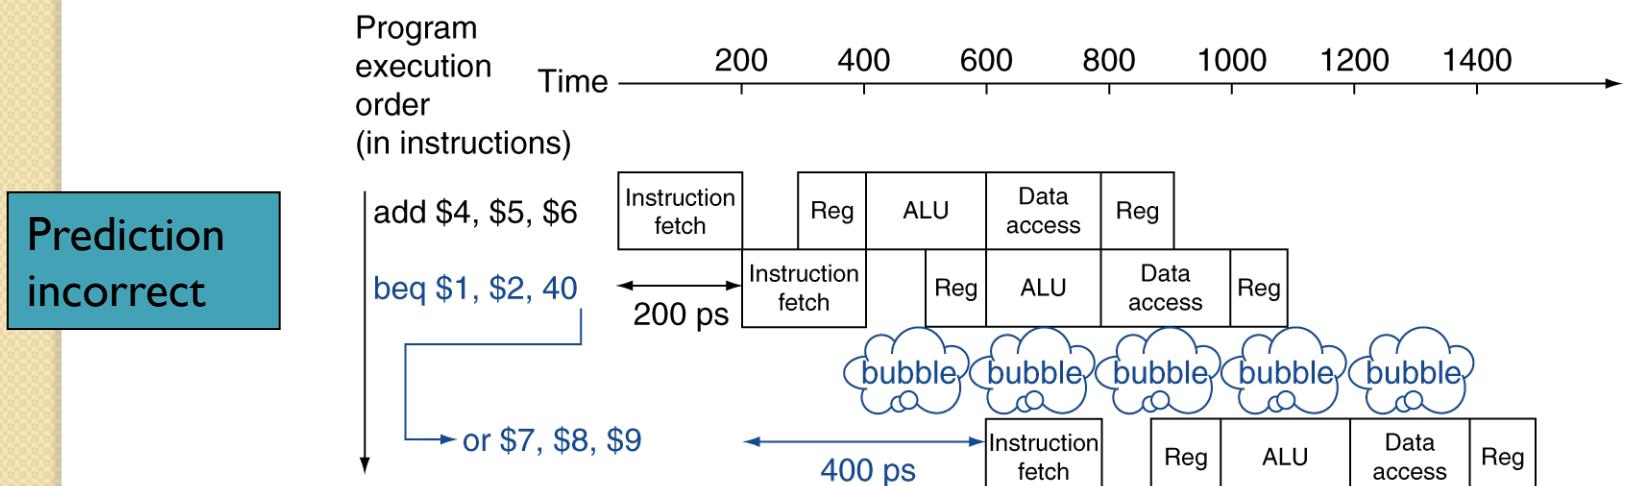
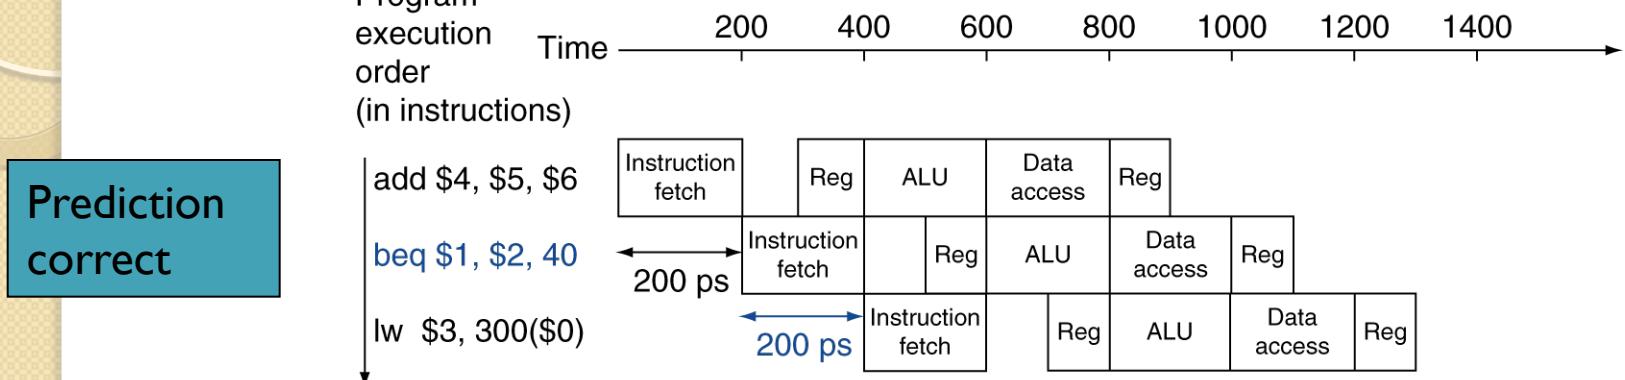
- Branch determines flow of control
 - Fetching next instruction depends on branch outcome
 - Pipeline can't always fetch correct instruction
 - Still working on ID stage of branch
- In MIPS pipeline
 - Need to compare registers and compute target early in the pipeline
 - Add hardware to do it in ID stage

Stall on Branch

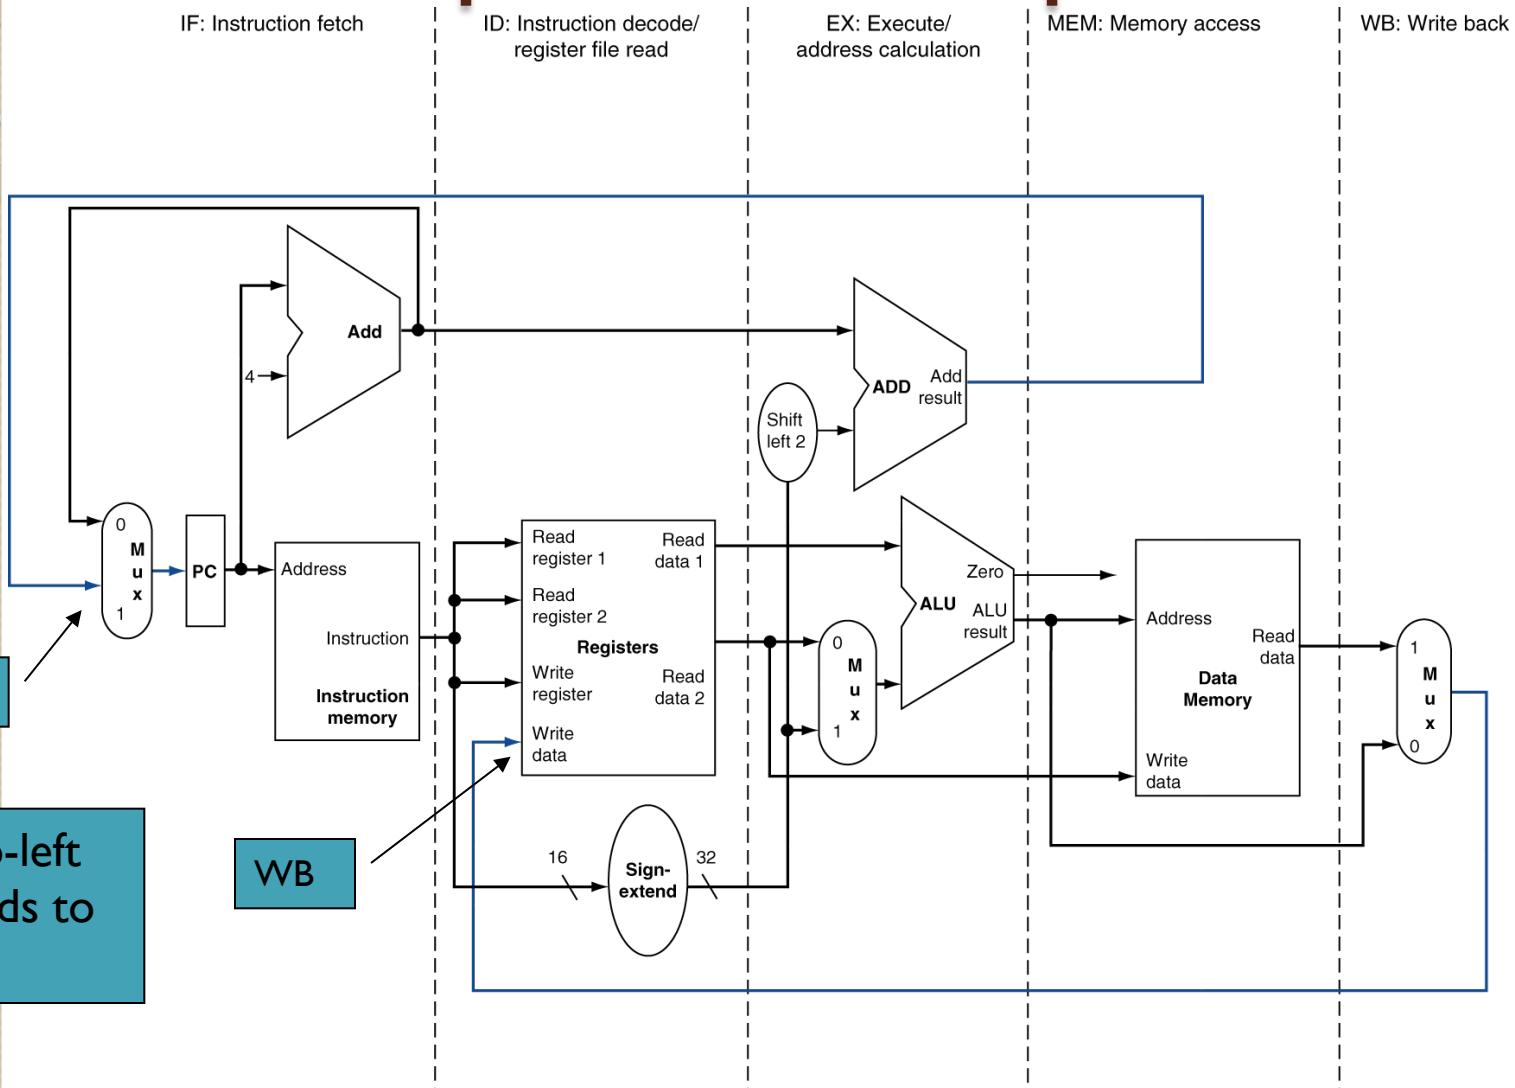
- Wait until branch outcome determined before fetching next instruction



MIPS with Predict Not Taken

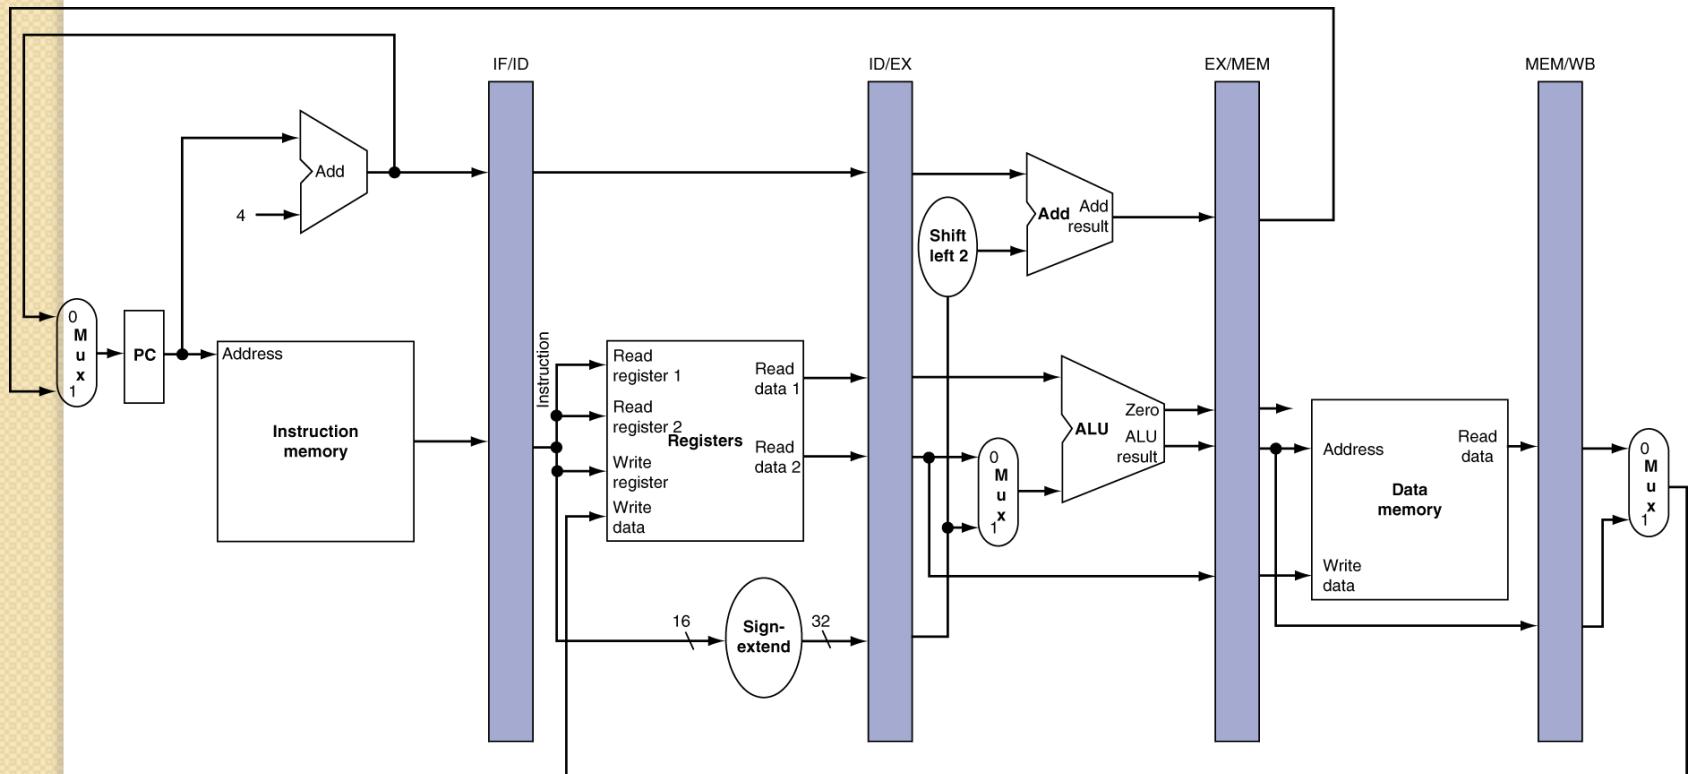


MIPS Pipelined Datapath



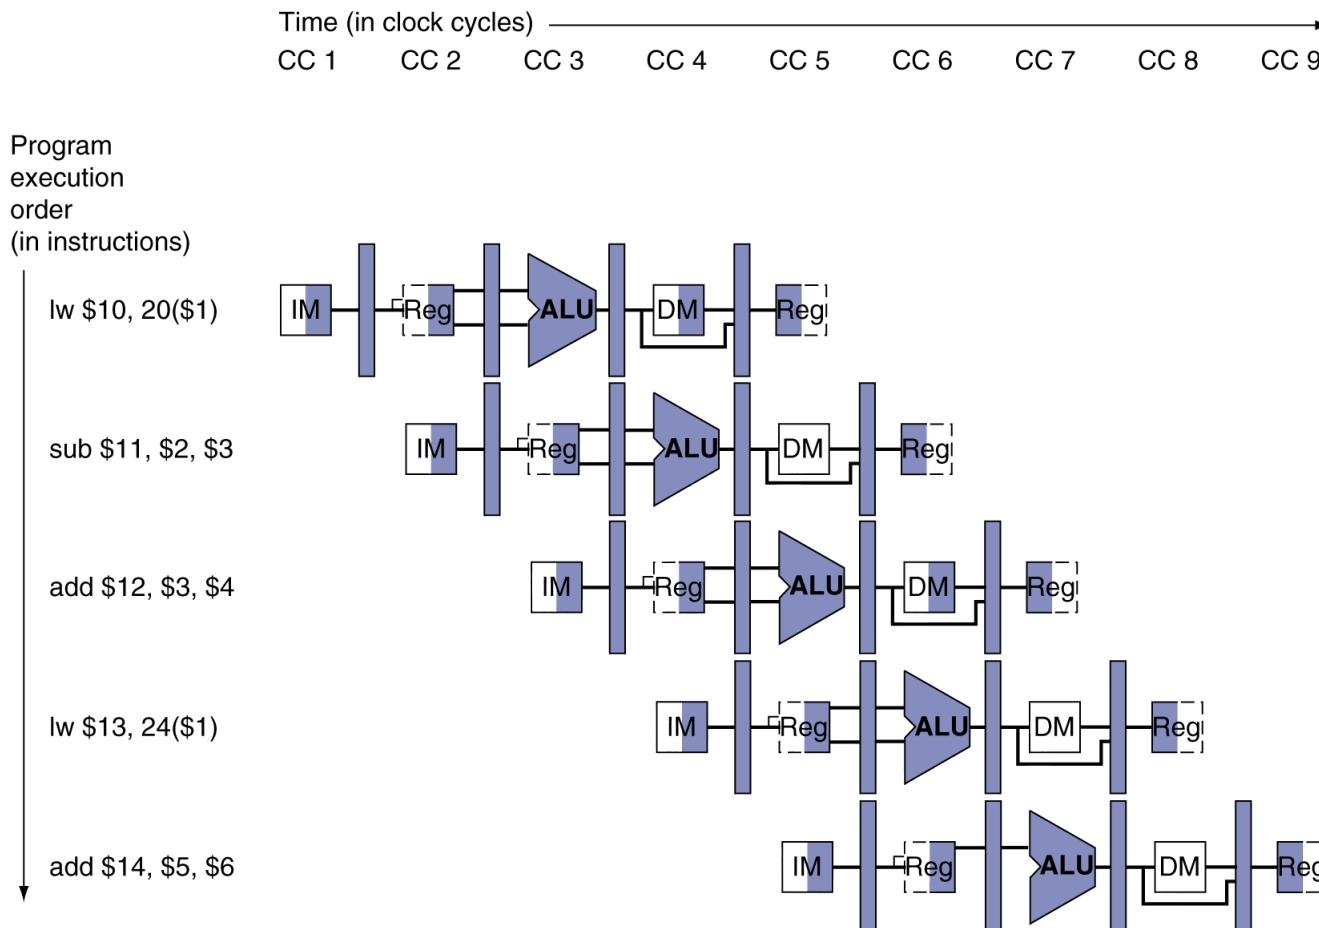
Pipeline registers

- Need registers between stages (Why ?)



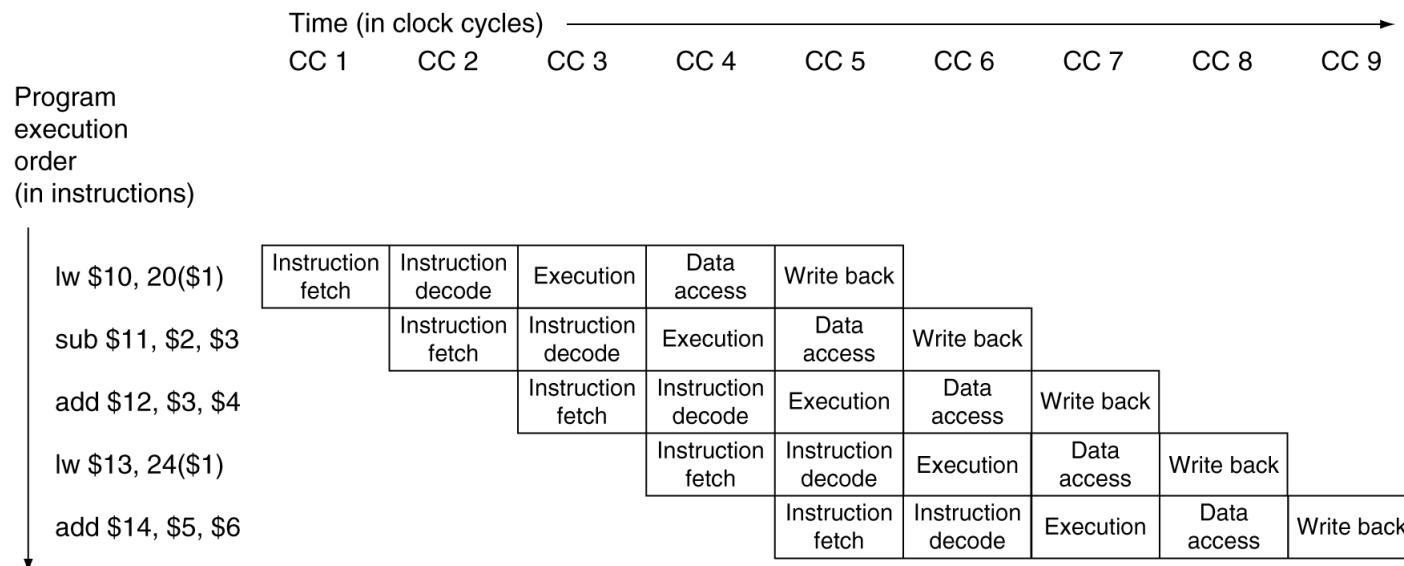
Multi-Cycle Pipeline Diagram

- Form showing resource usage



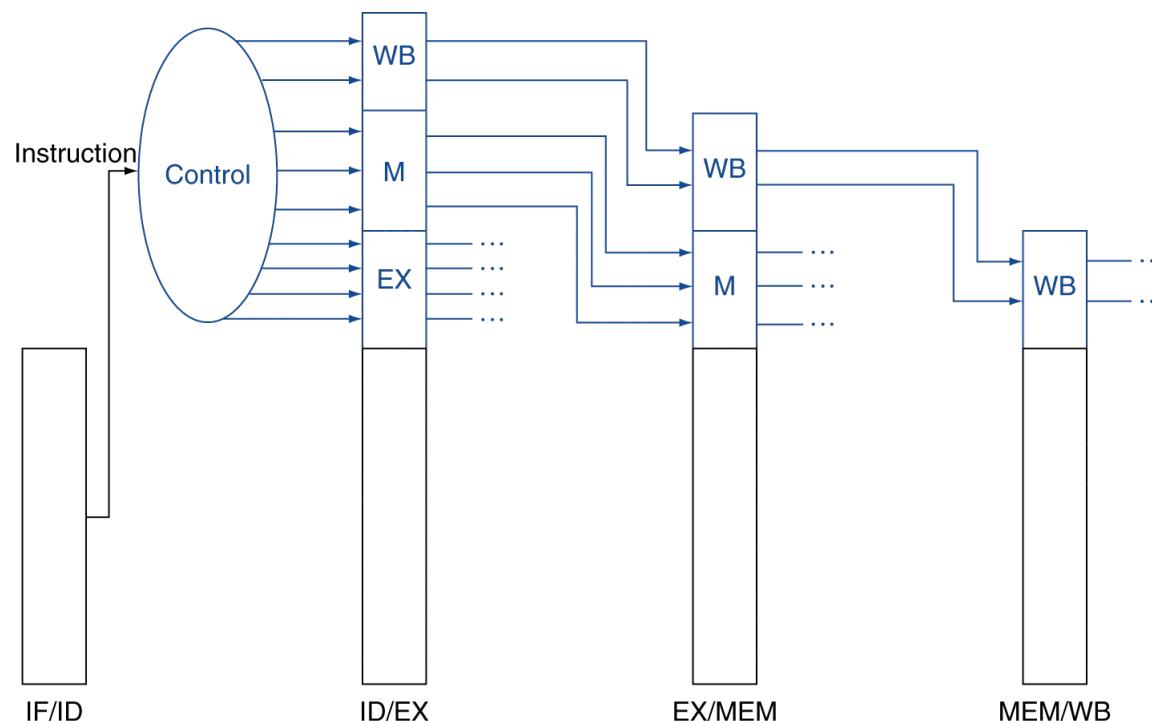
Multi-Cycle Pipeline Diagram

Traditional form



Pipelined Control

- Control signals derived from instruction
 - As in single-cycle implementation



Instruction-Level Parallelism (ILP)

- Pipelining: executing multiple instructions in parallel
- To increase ILP
 - Deeper pipeline
 - Less work per stage \Rightarrow shorter clock cycle
 - Multiple issue
 - Replicate pipeline stages \Rightarrow multiple pipelines
 - Start multiple instructions per clock cycle
 - CPI < 1, so use Instructions Per Cycle (IPC)
 - E.g., 4GHz 4-way multiple-issue
 - 16 BIPS, peak CPI = 0.25, peak IPC = 4
 - But dependencies reduce this in practice

MIPS with Static Dual Issue

- Two-issue packets
 - One ALU/branch instruction
 - One load/store instruction
 - 64-bit aligned
 - ALU/branch, then load/store
 - Pad an unused instruction with nop

Address	Instruction type	Pipeline Stages						
n	ALU/branch	IF	ID	EX	MEM	WB		
n + 4	Load/store	IF	ID	EX	MEM	WB		
n + 8	ALU/branch		IF	ID	EX	MEM	WB	
n + 12	Load/store		IF	ID	EX	MEM	WB	
n + 16	ALU/branch			IF	ID	EX	MEM	WB
n + 20	Load/store			IF	ID	EX	MEM	WB

Loop Unrolling Example

	ALU/branch	Load/store	cycle
Loop:	addi \$s1, \$s1,-16	lw \$t0, 0(\$s1)	1
	nop	lw \$t1, 12(\$s1)	2
	addu \$t0, \$t0, \$s2	lw \$t2, 8(\$s1)	3
	addu \$t1, \$t1, \$s2	lw \$t3, 4(\$s1)	4
	addu \$t2, \$t2, \$s2	sw \$t0, 16(\$s1)	5
	addu \$t3, \$t4, \$s2	sw \$t1, 12(\$s1)	6
	nop	sw \$t2, 8(\$s1)	7
	bne \$s1, \$zero, Loop	sw \$t3, 4(\$s1)	8

- $\text{IPC} = 14/8 = 1.75$
 - Closer to 2, but at cost of registers and code size

Dynamic Multiple Issue

- “Superscalar” processors
- CPU decides whether to issue 0, 1, 2, ... each cycle
 - Avoiding structural and data hazards
- Avoids the need for compiler scheduling
 - Though it may still help
 - Code semantics ensured by the CPU

Does Multiple Issue Work?

The BIG Picture

- Yes, but not as much as we'd like
- Programs have real dependencies that limit ILP
- Some dependencies are hard to eliminate
 - e.g., pointer aliasing
- Some parallelism is hard to expose
 - Limited window size during instruction issue
- Memory delays and limited bandwidth
 - Hard to keep pipelines full
- Speculation can help if done well

Concluding Remarks for Processors

- ISA influences design of datapath and control
- Datapath and control influence design of ISA
- Pipelining improves instruction throughput using parallelism
 - More instructions completed per second
 - Latency for each instruction not reduced
- Hazards: structural, data, control
- Multiple issue and dynamic scheduling (ILP)
 - Dependencies limit achievable parallelism
 - Complexity leads to the power wall



QUESTIONS



**What is the difference between a
conflict miss and a compulsory miss?
How would you reduce each type?**

What are two different strategies for dealing with cache writes? What is an advantage and disadvantage of each type?

- **Why might we want more than one level of a cache?**

- Suppose a processor has a **CPI** of 3.0 given a perfect cache. If there are 1.4 memory accesses per instruction, a miss penalty of 20 cycles, and a miss rate of 5%, what is the effective **CPI** with the real cache?

- **What are two advantages of using virtual memory ?**

- **What is a TLB? Why do we need it?**
- **Design a cache controller.**

- **How does pipelining improve performance?**

- **What is multiple issue? Is this the same as VLIW?**

Why is branch prediction so important? Does multiple issue increase or decrease the need for such prediction?

- Explain the stored-program concept.

- For a data cache with a 92% hit rate and a 2-cycle hit latency, calculate the average memory access latency (AMAT). Assume that the latency to memory and a cache miss penalty together is 124 clock cycles.
 - $$\begin{aligned} \text{AMAT} &= 2 + 0.08 * 124 \\ &= 11.92 \text{ cycles} \end{aligned}$$

Yes, CPI ?

(5 pts) A compiler designer is trying to decide between two code segments for a particular machine. The hardware designers have provided the following data below about the CPI for each class, and the instruction counts being considered for each code sequence.

Class	CPI for this instruction class
A	2
B	3

	Instruction Counts for Instruction Classes	
Code sequence	A	B
1	3	5
2	7	2

How many cycles are required for each code sequence?

Code sequence #1:

Code sequence #2:

Which is faster and how by how much?

What is the CPI for each code sequence?

CPI for code sequence #1:

CPI for code sequence #2: